

kurs


December 26, 2022


<https://paperswithcode.com/dataset/oxford-102-flower>

1 Oxford 102 Flower

oxford 102 flower , 102 ,
40 258 .
 , , .
 , .

1.1





1.2

paperswithcode.com

:

-

-

-

: - 102

(RGB

) +

-

,

.

1.3

Big Transfer (BiT): General Visual Representation Learning (2019)

Big Transfer :

,


.

,

.

BiT-L, SOTA, ResNet-50

.



to downstream tasks - : - Upstream Pre-train - - Transfer

1.4

1.4.1

```
[167]: import tensorflow as tf
import tensorflow_hub as hub

import time

from PIL import Image, ImageStat
import requests
from io import BytesIO

import matplotlib.pyplot as plt
import numpy as np

import glob, os
import pathlib

import plotly as xplt
import plotly.express as px
import plotly.graph_objects as go
import datetime
```

```
[98]: tf_flowers_labels = ['dandelion', 'daisy', 'tulips', 'sunflowers', 'roses']
```

ResNet50

```
[3]: model_url = "https://tfhub.dev/google/bit/m-r50x1/1"
      module = hub.KerasLayer(model_url)
```

Metal device set to: Apple M1 Pro

```
2022-12-26 14:01:16.990700: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305]
Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel
may not have been built with NUMA support.
2022-12-26 14:01:16.990917: I
tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271]
Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0
MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id:
<undefined>)
```

```
[4]: data_dir = tf.keras.utils.get_file(origin='https://storage.googleapis.com/
      ↪download.tensorflow.org/example_images/flower_photos.tgz',
      fname='flower_photos', untar=True)
      data_dir = pathlib.Path(data_dir)
```

```
[176]: from plotly.subplots import make_subplots
import plotly.graph_objects as go
from statistics import mean
import cv2

def plotdistribhw(imgs):
    df = {'width': [], 'height': []}
    for img in imgs:
        image = Image.open(img)
        df['width'].append(image.size[0])
        df['height'].append(image.size[1])
    fig = make_subplots(rows=1, cols=2)
    fig.add_trace(go.Histogram(x = df['width'], name = 'width',1,1))
    fig.add_trace(go.Histogram(x = df['height'], name = 'height',1,2))
    fig.update_layout(title_text='Histogram of image dimensions')
    fig.show()

def plotdistribcolor(imgs):
    df = {'R': [], 'G': [], 'B': []}
    for img in imgs:
        image = Image.open(img)
        stat = ImageStat.Stat(image)
        df['R'].append(stat.mean[0])
        df['G'].append(stat.mean[1])
        df['B'].append(stat.mean[2])
```

```

fig = go.Figure()
fig.add_trace(go.Histogram(x = df['R'], name = 'Red Channel'))
fig.add_trace(go.Histogram(x = df['G'], name = 'Green Channel'))
fig.add_trace(go.Histogram(x = df['B'], name = 'Blue Channel'))
fig.update_layout(title_text=' ')
fig.show()

def datasetinfo(data_dir, tf_flowers_labels):
    imgs = glob.glob(f'{data_dir}/**/*.jpg', recursive=True)

    fig = make_subplots(rows=1, cols=5)
    for i, img in enumerate(imgs[55:60]):
        x = Image.open(img)
        fig.add_trace(go.Image(z=x), 1, i+1)
    fig.update_layout(
        height=300,
        title_text='Picture examples from dataset'
    )
    fig.show()

    print(f' : {len(imgs)}\n ( ) : \
↪ {len(tf_flowers_labels)}')

    plotdistribhw(imgs)

    plotdistribcolor(imgs)

```


1.4.2

```
[71]: print(data_dir)
```

/Users/ivanskvortsov/.keras/datasets/flower_photos

```
[177]: datasetinfo(data_dir, tf_flowers_labels)
```

Picture examples from dataset




```

row = f.readline()
row = row.rstrip()
imagenet_int_to_str.update({i: row})

```

```

--2022-12-26 14:01:48--
https://storage.googleapis.com/bit_models/ilsvrc2012_wordnet_lemmas.txt
Resolving storage.googleapis.com (storage.googleapis.com)... 74.125.205.128,
64.233.165.128, 173.194.73.128, ...
Connecting to storage.googleapis.com
(storage.googleapis.com)|74.125.205.128|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 21675 (21K) [text/plain]
Saving to: 'ilsvrc2012_wordnet_lemmas.txt.3'

ilsvrc2012_wordnet_ 100%[=====>] 21,17K --.-KB/s in 0,02s

2022-12-26 14:01:48 (1,04 MB/s) - 'ilsvrc2012_wordnet_lemmas.txt.3' saved
[21675/21675]

```

1.5

-

```

[9]: IMG_HEIGHT = 224
    IMG_WIDTH = 224

    CLASS_NAMES = tf_flowers_labels # from plotting helper functions above
    NUM_CLASSES = len(CLASS_NAMES)
    num_examples = len(list(data_dir.glob('*/*.jpg')))

    def get_label(file_path):
        # convert the path to a list of path components
        parts = tf.strings.split(file_path, os.path.sep)
        # The second to last is the class-directory

        return tf.where(parts[-2] == CLASS_NAMES)[0][0]

    def decode_img(img):
        # convert the compressed string to a 3D uint8 tensor
        img = tf.image.decode_jpeg(img, channels=3)
        return img

    def process_path(file_path):
        label = get_label(file_path)
        # load the raw data from the file as a string
        img = tf.io.read_file(file_path)
        img = decode_img(img)

```



```

features = {'image': img, 'label': label}
return features

list_ds = tf.data.Dataset.list_files(str(data_dir/'*/'))
ds = list_ds.map(process_path, num_parallel_calls=tf.data.experimental.AUTOTUNE)

```

```

[10]: def preprocess_image(image):
    image = np.array(image)
    # reshape into shape [batch_size, height, width, num_channels]
    img_resized = tf.reshape(image, [1, image.shape[0], image.shape[1], image.
↪shape[2]])
    # Use `convert_image_dtype` to convert to floats in the [0,1] range.
    image = tf.image.convert_image_dtype(img_resized, tf.float32)
    return image

def load_image_from_url(url):
    """Returns an image with shape [1, height, width, num_channels]."""
    response = requests.get(url)
    image = Image.open(BytesIO(response.content))
    image = preprocess_image(image)
    return image

```

1.5.1

```

[11]: # Show the MAX_PREDS highest scoring labels:
MAX_PREDS = 5
# Do not show labels with lower score than this:
MIN_SCORE = 0.8

def show_preds(logits, image, correct_flowers_label=None, ↵
↪tf_flowers_logits=False):

    if len(logits.shape) > 1:
        logits = tf.reshape(logits, [-1])

    fig, axes = plt.subplots(1, 2, figsize=(7, 4), squeeze=False)

    ax1, ax2 = axes[0]

    ax1.axis('off')
    ax1.imshow(image)
    if correct_flowers_label is not None:
        ax1.set_title(tf_flowers_labels[correct_flowers_label])
    classes = []
    scores = []

```

```

logits_max = np.max(logits)
softmax_denominator = np.sum(np.exp(logits - logits_max))
for index, j in enumerate(np.argsort(logits)[-MAX_PREDS:][::-1]):
    score = 1.0/(1.0 + np.exp(-logits[j]))
    if score < MIN_SCORE: break
    if not tf_flowers_logits:
        # predicting in imagenet label space
        classes.append(imagenet_int_to_str[j])
    else:
        # predicting in tf_flowers label space
        classes.append(tf_flowers_labels[j])
    scores.append(np.exp(logits[j] - logits_max)/softmax_denominator*100)

ax2.barh(np.arange(len(scores)) + 0.1, scores)
ax2.set_xlim(0, 100)
ax2.set_yticks(np.arange(len(scores)))
ax2.yaxis.set_ticks_position('right')
ax2.set_yticklabels(classes, rotation=0, fontsize=14)
ax2.invert_xaxis()
ax2.invert_yaxis()
ax2.set_xlabel('Prediction probabilities', fontsize=11)

```

```
[12]: print(data_dir)
```

```
/Users/ivanskvortsov/.keras/datasets/flower_photos
```

1.6 Showcase

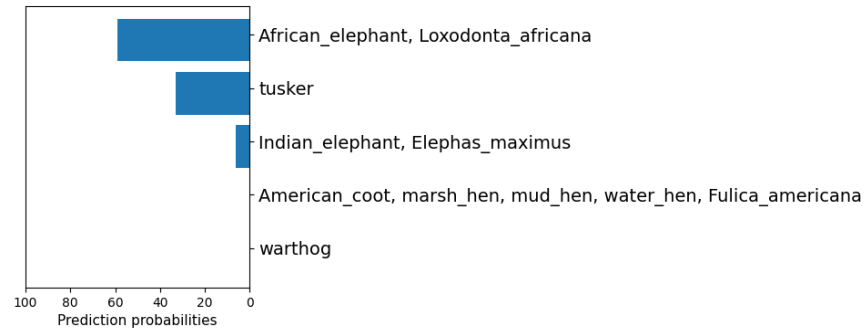
```
[13]: model_url = "https://tfhub.dev/google/bit/m-r50x1/ilsrvrc2012_classification/1"
imagenet_module = hub.KerasLayer(model_url)
```

```
[16]: # Load image (image provided is CC0 licensed)
img_url = "http://images6.fanpop.com/image/photos/34700000/
↳Grey-Elephant-colors-34712059-620-388.jpg"
image = load_image_from_url(img_url)

# Run model on image
logits = imagenet_module(image)

# Show image and predictions
show_preds(logits, image[0])

```



```
[17]: train_split = 0.9
num_train = int(train_split * num_examples)
ds_train = ds.take(num_train)
ds_test = ds.skip(num_train)


DATASET_NUM_TRAIN_EXAMPLES = num_examples
```

```
[18]: for features in ds_train.take(1):
    image = features['image']
    image = preprocess_image(image)

    # Run model on image
    logits = imagenet_module(image)

    # Show image and predictions
    show_preds(logits, image[0], correct_flowers_label=features['label'].numpy())
```

2022-12-26 14:02:57.453063: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.



1.7 ()

1.7.1

```
[19]: # Add new head to the BiT model

class MyBiTModel(tf.keras.Model):
    """BiT with a new head."""

    def __init__(self, num_classes, module):
        super().__init__()

        self.num_classes = num_classes
        self.head = tf.keras.layers.Dense(num_classes, kernel_initializer='zeros')
        self.bit_model = module

    def call(self, images):
        # No need to cut head off since we are using feature extractor model
        bit_embedding = self.bit_model(images)
        return self.head(bit_embedding)

model = MyBiTModel(num_classes=NUM_CLASSES, module=module)
```

```
[20]: IMAGE_SIZE = "\u003C96x96 px" #@param ["<96x96 px", "> 96 x 96 px"]
DATASET_SIZE = "\u003C20k examples" #@param ["<20k examples", "20k-500k",
    ↪ "examples", ">500k examples"]

if IMAGE_SIZE == "<96x96 px":
    RESIZE_TO = 160
    CROP_TO = 128
else:
```

```

RESIZE_TO = 512
CROP_TO = 480

if DATASET_SIZE == "<20k examples":
    SCHEDULE_LENGTH = 500
    SCHEDULE_BOUNDARIES = [200, 300, 400]
elif DATASET_SIZE == "20k-500k examples":
    SCHEDULE_LENGTH = 10000
    SCHEDULE_BOUNDARIES = [3000, 6000, 9000]
else:
    SCHEDULE_LENGTH = 20000
    SCHEDULE_BOUNDARIES = [6000, 12000, 18000]

```

[21]: *# Preprocessing helper functions*

```

# Create data pipelines for training and testing:
BATCH_SIZE = 512
SCHEDULE_LENGTH = SCHEDULE_LENGTH * 512 / BATCH_SIZE

STEPS_PER_EPOCH = 10

def cast_to_tuple(features):
    return (features['image'], features['label'])

def preprocess_train(features):
    # Apply random crops and horizontal flips for all tasks
    # except those for which cropping or flipping destroys the label semantics
    # (e.g. predict orientation of an object)
    features['image'] = tf.image.random_flip_left_right(features['image'])
    features['image'] = tf.image.resize(features['image'], [RESIZE_TO, RESIZE_TO])
    features['image'] = tf.image.random_crop(features['image'], [CROP_TO,
↪CROP_TO, 3])
    features['image'] = tf.cast(features['image'], tf.float32) / 255.0
    return features

def preprocess_test(features):
    features['image'] = tf.image.resize(features['image'], [RESIZE_TO, RESIZE_TO])
    features['image'] = tf.cast(features['image'], tf.float32) / 255.0
    return features

pipeline_train = (ds_train
                  .shuffle(10000)
                  .repeat(int(SCHEDULE_LENGTH * BATCH_SIZE /
↪DATASET_NUM_TRAIN_EXAMPLES * STEPS_PER_EPOCH) + 1 + 50) # repeat
↪dataset_size / num_steps
                  .map(preprocess_train, num_parallel_calls=8)
                  .batch(BATCH_SIZE))

```

```

        .map(cast_to_tuple) # for keras model.fit
        .prefetch(2))

pipeline_test = (ds_test.map(preprocess_test, num_parallel_calls=1)
                  .map(cast_to_tuple) # for keras model.fit
                  .batch(BATCH_SIZE)
                  .prefetch(2))

```

```

[22]: # Define optimiser and loss

lr = 0.003 * BATCH_SIZE / 512

# Decay learning rate by a factor of 10 at SCHEDULE_BOUNDARIES.
lr_schedule = tf.keras.optimizers.schedules.
    ↳ PiecewiseConstantDecay(boundaries=SCHEDULE_BOUNDARIES,
                            values=[lr,
    ↳ lr*0.1, lr*0.001, lr*0.0001])
optimizer = tf.keras.optimizers.SGD(learning_rate=lr_schedule, momentum=0.9)

loss_fn = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True)

```

1.7.2

```

[27]: %load_ext tensorboard

```

```

[32]: # model.compile(optimizer=optimizer,
#                   loss=loss_fn,
#                   metrics=['accuracy'])

# log_dir = "./logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
# tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir)#,
    ↳ histogram_freq=1)

# # histogram_freq=1 - ( )

# # Fine-tune model
# history = model.fit(
#     pipeline_train,
#     batch_size=BATCH_SIZE,
#     steps_per_epoch=STEPS_PER_EPOCH,
#     epochs=10,
#     validation_data=pipeline_test,
#     callbacks=[tensorboard_callback]
# )

```

Epoch 1/10

```

2022-12-26 14:40:56.799844: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

10/10 [=====] - ETA: 0s - loss: 0.0626 - accuracy:
0.9812

2022-12-26 14:41:36.350812: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.

WARNING:tensorflow:6 out of the last 6 calls to <function
Model.make_test_function.<locals>.test_function at 0x2e6ce8ee0> triggered
tf.function retracing. Tracing is expensive and the excessive number of tracings
could be due to (1) creating @tf.function repeatedly in a loop, (2) passing
tensors with different shapes, (3) passing Python objects instead of tensors.
For (1), please define your @tf.function outside of the loop. For (2),
@tf.function has reduce_retracing=True option that can avoid unnecessary
retracing. For (3), please refer to
https://www.tensorflow.org/guide/function#controlling\_retracing and
https://www.tensorflow.org/api\_docs/python/tf/function for more details.

WARNING:tensorflow:6 out of the last 6 calls to <function
Model.make_test_function.<locals>.test_function at 0x2e6ce8ee0> triggered
tf.function retracing. Tracing is expensive and the excessive number of tracings
could be due to (1) creating @tf.function repeatedly in a loop, (2) passing
tensors with different shapes, (3) passing Python objects instead of tensors.
For (1), please define your @tf.function outside of the loop. For (2),
@tf.function has reduce_retracing=True option that can avoid unnecessary
retracing. For (3), please refer to
https://www.tensorflow.org/guide/function#controlling\_retracing and
https://www.tensorflow.org/api\_docs/python/tf/function for more details.

10/10 [=====] - 59s 5s/step - loss: 0.0626 - accuracy:
0.9812 - val_loss: 0.0389 - val_accuracy: 0.9891
Epoch 2/10
10/10 [=====] - 25s 3s/step - loss: 0.0599 - accuracy:
0.9826 - val_loss: 0.0699 - val_accuracy: 0.9782
Epoch 3/10
10/10 [=====] - 26s 3s/step - loss: 0.0542 - accuracy:
0.9822 - val_loss: 0.0503 - val_accuracy: 0.9891
Epoch 4/10
10/10 [=====] - 24s 2s/step - loss: 0.0512 - accuracy:
0.9863 - val_loss: 0.0231 - val_accuracy: 0.9973
Epoch 5/10
10/10 [=====] - 26s 3s/step - loss: 0.0441 - accuracy:
0.9852 - val_loss: 0.0296 - val_accuracy: 0.9864
Epoch 6/10
10/10 [=====] - 23s 2s/step - loss: 0.0424 - accuracy:
0.9863 - val_loss: 0.0431 - val_accuracy: 0.9837

```

```
Epoch 7/10
10/10 [=====] - 24s 2s/step - loss: 0.0371 - accuracy:
0.9873 - val_loss: 0.0158 - val_accuracy: 1.0000
Epoch 8/10
10/10 [=====] - 26s 3s/step - loss: 0.0369 - accuracy:
0.9904 - val_loss: 0.0278 - val_accuracy: 0.9946
Epoch 9/10
10/10 [=====] - 23s 2s/step - loss: 0.0391 - accuracy:
0.9875 - val_loss: 0.0458 - val_accuracy: 0.9864
Epoch 10/10
10/10 [=====] - 22s 2s/step - loss: 0.0317 - accuracy:
0.9906 - val_loss: 0.0469 - val_accuracy: 0.9864
```

```
[33]: %tensorboard --logdir logs/fit
```

```
<IPython.core.display.HTML object>
```

```
[36]: # #Save fine-tuned model as SavedModel
# export_module_dir = './tmp/my_saved_bit_model/'
# tf.saved_model.save(model, export_module_dir)
```

```
INFO:tensorflow:Assets written to: ./tmp/my_saved_bit_model/assets
```

```
INFO:tensorflow:Assets written to: ./tmp/my_saved_bit_model/assets
```

1.7.3 showcase

```
[37]: saved_module = hub.KerasLayer(export_module_dir, trainable=True)
```

```
2022-12-26 14:49:02.825135: W
tensorflow/core/common_runtime/graph_constructor.cc:805] Node
're_lu_48/PartitionedCall' has 1 outputs but the _output_shapes attribute
specifies shapes for 2 outputs. Output shapes may be inaccurate.
2022-12-26 14:49:02.825908: W
tensorflow/core/common_runtime/graph_constructor.cc:805] Node
'global_average_pooling2d/PartitionedCall' has 1 outputs but the _output_shapes
attribute specifies shapes for 4 outputs. Output shapes may be inaccurate.
```

```
[38]: for features in ds_train.take(1):
    image = features['image']
    image = preprocess_image(image)
    image = tf.image.resize(image, [CROP_TO, CROP_TO])

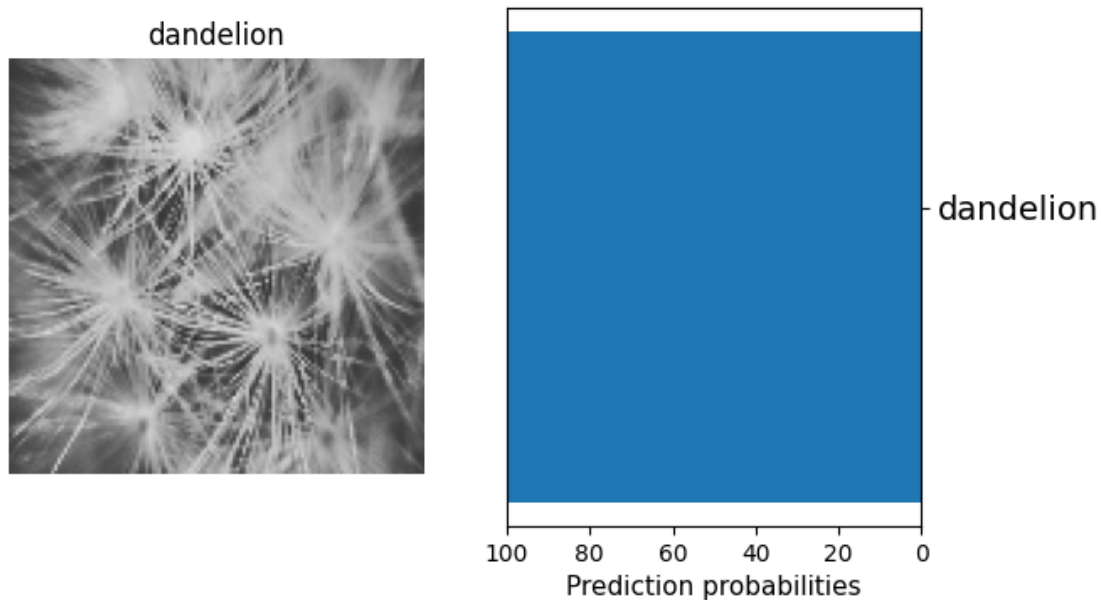
    # Run model on image
    logits = saved_module(image)

    # Show image and predictions
```



```
show_preds(logits, image[0], correct_flowers_label=features['label'].numpy(),  
↳tf_flowers_logits=True)
```

2022-12-26 14:49:07.590225: I
tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113]
Plugin optimizer for device_type GPU is enabled.



1.8 Preprocessing ESRGAN

Enhanced Super Resolution GAN

1.8.1 -

```
[180]: def preprocess_image(image_path):  
        """ Loads image from path and preprocesses to make it model ready  
        Args:  
            image_path: Path to the image file  
        """  
        hr_image = tf.image.decode_image(tf.io.read_file(image_path))  
        # If PNG, remove the alpha channel. The model only supports  
        # images with 3 color channels.  
        if hr_image.shape[-1] == 4:  
            hr_image = hr_image[..., :-1]  
        hr_size = (tf.convert_to_tensor(hr_image.shape[:-1]) // 4) * 4  
        hr_image = tf.image.crop_to_bounding_box(hr_image, 0, 0, hr_size[0],  
↳hr_size[1])
```

```

    hr_image = tf.cast(hr_image, tf.float32)
    return tf.expand_dims(hr_image, 0)

def save_image(image, filename):
    """
    Saves unscaled Tensor Images.
    Args:
        image: 3D image tensor. [height, width, channels]
        filename: Name of the file to save.
    """
    if not isinstance(image, Image.Image):
        image = tf.clip_by_value(image, 0, 255)
        image = Image.fromarray(tf.cast(image, tf.uint8).numpy())
    image.save("%s.jpg" % filename)
    print("Saved as %s.jpg" % filename)

def plot_image(image, title=""):
    """
    Plots images from image tensors.
    Args:
        image: 3D image tensor. [height, width, channels].
        title: Title to display in the plot.
    """
    image = np.asarray(image)
    image = tf.clip_by_value(image, 0, 255)
    image = Image.fromarray(tf.cast(image, tf.uint8).numpy())
    plt.imshow(image)
    plt.axis("off")
    plt.title(title)

```

1.8.2 ESRGAN

```

[181]: ESR_SAVED_MODEL_PATH = "https://tfhub.dev/captain-pool/esrgan-tf2/1"
       IMAGE_PATH="image.jpg"

```

ESRGAN tfhub

```

[182]: hr_image = preprocess_image(IMAGE_PATH)
       plot_image(tf.squeeze(hr_image), title="Original Image")
       save_image(tf.squeeze(hr_image), filename="Original Image")

```

Saved as Original Image.jpg

Original Image



```
[183]: model = hub.load(ESR_SAVED_MODEL_PATH)
```

ESRGAN,

```
[185]: start = time.time()
fake_image = model(hr_image)
fake_image = tf.squeeze(fake_image)
print("Time Taken: %f" % (time.time() - start))
plot_image(tf.squeeze(fake_image), title="Super Resolution")
save_image(tf.squeeze(fake_image), filename="Super Resolution")
```

Time Taken: 0.193167

Saved as Super Resolution.jpg

Super Resolution



1.9

ESRGAN, oxford102. ResNet50, TensorBoard. Plotly.

```
[186]: !export PATH=/Library/TeX/texbin:$PATH
```

```
[ ]:
```