

<https://paperswithcode.com/dataset/oxford-102-flower>

# Oxford 102 Flower

**oxford 102 flower** это датасет составленный для задач классификации, состоящий из 102 категорий цветов, в основном встречающихся на территории Великобритании. Каждый класс состоит от 40 до 258 картинок класса.

Картинки имеют достаточно большое разрешение, различные ракурсы, цветовые палитры и освещение. Плюс к этому есть категории которые имеют сильно отличающиеся друг от друга картинки и несколько похожих категорий.

Изначально датасет был сформирован для задачи классификации по большому количеству классов, которые достаточно похожи.

## Визуализации

## Решаемые задачи

Согласно *paperswithcode.com* датасет используется для следующих задач:

- Классификация изображений
- Генерация картинки по тексту
- Генерации картинок
- Инкрементальное обучение

Датасет представлен:

- 102 категории цветов (RGB картинки) + лейблы
- Маски сегментации

Однако маски нам не нужны, так что будем использовать версию датасета без них.


## Статья

[Big Transfer \(BiT\): General Visual Representation Learning](#) (2019)

В статье предложено решение проблемы эффективности тренировки нейронных сетей на собственных семплах и упрощение гиперпараметров при тренировке посредством операции **Big Transfer** : Суть метода в тренировке нейронной сети на

большом, хорошо изученном и модерированном датасете. Авторы, применяя сильно эскалированный претренин и перенос параметров достигают хорошей точности на большом количестве датасетов.

На картинке видим как перенос перформанса с BiT-L, SOTA, ResNet-50 на конкретные задачи влияет на точность.

 Screenshot 2022-12-26 at 19.46.29.png

Сама методика состоит из нескольких частей:

- Upstream Pre-train - тренировка на очень большом датасете изначальной модели
- Transfer to downstream tasks - выбираем только важные гиперпараметры и тренируем модель на данных для задачи с большим количеством препроцессинга - обрез картинок, перевероты.

## Импорт и загрузка датасета

### Импорты и загрузка

```
In [167]: import tensorflow as tf
import tensorflow_hub as hub

import time

from PIL import Image, ImageStat
import requests
from io import BytesIO

import matplotlib.pyplot as plt
import numpy as np

import glob, os
import pathlib

import plotly as xplt
import plotly.express as px
import plotly.graph_objects as go
import datetime
```

Определим лейблы для переноса

```
In [98]: tf_flowers_labels = ['dandelion', 'daisy', 'tulips', 'sunflowers', 'roses']
```

Возьмем готовые веса модели ResNet50

```
In [3]: model_url = "https://tfhub.dev/google/bit/m-r50x1/1"
module = hub.KerasLayer(model_url)
```

Metal device set to: Apple M1 Pro

```
2022-12-26 14:01:16.990700: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305] Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel may not have been built with NUMA support.
2022-12-26 14:01:16.990917: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271] Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0 MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id: <undefined>)
```

Скачаем и разархивируем датасет

```
In [4]: data_dir = tf.keras.utils.get_file(origin='https://storage.googleapis.com/d
                                             fname='flower_photos', untar=True)
data_dir = pathlib.Path(data_dir)
```

## Хелпер-функции для датасет инфо

```
In [176... from plotly.subplots import make_subplots
import plotly.graph_objects as go
from statistics import mean
import cv2

def plotdistribhw(imgs):
    df = {'width': [], 'height': []}
    for img in imgs:
        image = Image.open(img)
        df['width'].append(image.size[0])
        df['height'].append(image.size[1])
    fig = make_subplots(rows=1, cols=2)
    fig.add_trace(go.Histogram(x = df['width'], name = 'Распределение по шир
    fig.add_trace(go.Histogram(x = df['height'], name = 'Распределение по в
    fig.update_layout(title_text='Размерности картинок')
    fig.show()

def plotdistribcolor(imgs):
    df = {'R': [], 'G': [], 'B': []}
    for img in imgs:
        image = Image.open(img)
        stat = ImageStat.Stat(image)
        df['R'].append(stat.mean[0])
        df['G'].append(stat.mean[1])
        df['B'].append(stat.mean[2])
    fig = go.Figure()
    fig.add_trace(go.Histogram(x = df['R'], name = 'Red Channel'))
    fig.add_trace(go.Histogram(x = df['G'], name = 'Green Channel'))
    fig.add_trace(go.Histogram(x = df['B'], name = 'Blue Channel'))
    fig.update_layout(title_text='Распределение по цветовым каналам')
    fig.show()

def datasetinfo(data_dir, tf_flowers_labels):
    imgs = glob.glob(f'{data_dir}/**/*.jpg', recursive=True)

    fig = make_subplots(rows=1, cols=5)
    for i, img in enumerate(imgs[55:60]):
```

```
x = Image.open(img)
fig.add_trace(go.Image(z=x), 1, i+1)
fig.update_layout(
    height=300,
    title_text='Picture examples from dataset'
)
fig.show()

print(f'Количество изображений: {len(imgs)}\nКоличество классов (урезанн

plotdistribhw(imgs)

plotdistribcolor(imgs)
```

## Информация о датасете

```
In [71]: print(data_dir)

/Users/ivanskvortsov/.keras/datasets/flower_photos
```

```
In [177... datasetinfo(data_dir, tf_flowers_labels)
```

Количество изображений: 3670  
Количество классов (урезанная версия): 5



## BiT-M R50x1

Feature extraction модель построенная на архитектуре ResNet50-v2, тренированная на большом количестве классов ImageNet-21k - датасет с 14 миллионами картинок и 21.843 классами. Выводом являются 2048-размерные фича векторы, модель используется для фэйн-тюнинга на новой задаче.

Ссылка: [ResNet](#)

Готовые леммы для ResNet50

```
In [8]: !wget https://storage.googleapis.com/bit_models/ilsvrc2012_wordnet_lemmas.tx

imagenet_int_to_str = {}

with open('ilsvrc2012_wordnet_lemmas.txt', 'r') as f:
    for i in range(1000):
        row = f.readline()
        row = row.rstrip()
        imagenet_int_to_str.update({i: row})
```

```
--2022-12-26 14:01:48-- https://storage.googleapis.com/bit_models/ilstvrc20
12_wordnet_lemmas.txt
Resolving storage.googleapis.com (storage.googleapis.com)... 74.125.205.12
8, 64.233.165.128, 173.194.73.128, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|74.125.205.12
8|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 21675 (21K) [text/plain]
Saving to: 'ilstvrc2012_wordnet_lemmas.txt.3'

ilstvrc2012_wordnet_ 100%[=====] 21,17K --.-KB/s in 0,02
s

2022-12-26 14:01:48 (1,04 MB/s) - 'ilstvrc2012_wordnet_lemmas.txt.3' saved
[21675/21675]
```

## Дообучение

Функции-хелперы для загрузки датасета и небольшой предобработки

```
In [9]: IMG_HEIGHT = 224
        IMG_WIDTH = 224

        CLASS_NAMES = tf_flowers_labels # from plotting helper functions above
        NUM_CLASSES = len(CLASS_NAMES)
        num_examples = len(list(data_dir.glob('*/*.jpg'))))

        def get_label(file_path):
            # convert the path to a list of path components
            parts = tf.strings.split(file_path, os.path.sep)
            # The second to last is the class-directory

            return tf.where(parts[-2] == CLASS_NAMES)[0][0]

        def decode_img(img):
            # convert the compressed string to a 3D uint8 tensor
            img = tf.image.decode_jpeg(img, channels=3)
            return img

        def process_path(file_path):
            label = get_label(file_path)
            # load the raw data from the file as a string
            img = tf.io.read_file(file_path)
            img = decode_img(img)
            features = {'image': img, 'label': label}
            return features

        list_ds = tf.data.Dataset.list_files(str(data_dir/'*/*'))
        ds = list_ds.map(process_path, num_parallel_calls=tf.data.experimental.AUTOTUNE)
```

Непосредственно препроцессинг

```
In [10]: def preprocess_image(image):
```

```

image = np.array(image)
# reshape into shape [batch_size, height, width, num_channels]
img_resized = tf.reshape(image, [1, image.shape[0], image.shape[1], image.shape[2]])
# Use `convert_image_dtype` to convert to floats in the [0,1] range.
image = tf.image.convert_image_dtype(img_resized, tf.float32)
return image

def load_image_from_url(url):
    """Returns an image with shape [1, height, width, num_channels]."""
    response = requests.get(url)
    image = Image.open(BytesIO(response.content))
    image = preprocess_image(image)
    return image

```

## Хелпер функции

```

In [11]: # Show the MAX_PREDS highest scoring labels:
MAX_PREDS = 5
# Do not show labels with lower score than this:
MIN_SCORE = 0.8

def show_preds(logits, image, correct_flowers_label=None, tf_flowers_logits=None):

    if len(logits.shape) > 1:
        logits = tf.reshape(logits, [-1])

    fig, axes = plt.subplots(1, 2, figsize=(7, 4), squeeze=False)

    ax1, ax2 = axes[0]

    ax1.axis('off')
    ax1.imshow(image)
    if correct_flowers_label is not None:
        ax1.set_title(tf_flowers_labels[correct_flowers_label])
    classes = []
    scores = []
    logits_max = np.max(logits)
    softmax_denominator = np.sum(np.exp(logits - logits_max))
    for index, j in enumerate(np.argsort(logits)[-MAX_PREDS:][::-1]):
        score = 1.0/(1.0 + np.exp(-logits[j]))
        if score < MIN_SCORE: break
        if not tf_flowers_logits:
            # predicting in imagenet label space
            classes.append(imagenet_int_to_str[j])
        else:
            # predicting in tf_flowers label space
            classes.append(tf_flowers_labels[j])
        scores.append(np.exp(logits[j] - logits_max)/softmax_denominator*100)

    ax2.barh(np.arange(len(scores)) + 0.1, scores)
    ax2.set_xlim(0, 100)
    ax2.set_yticks(np.arange(len(scores)))
    ax2.yaxis.set_ticks_position('right')
    ax2.set_yticklabels(classes, rotation=0, fontsize=14)
    ax2.invert_xaxis()

```



```
ax2.invert_yaxis()
ax2.set_xlabel('Prediction probabilities', fontsize=11)
```

```
In [12]: print(data_dir)

/Users/ivanskvortsov/.keras/datasets/flower_photos
```

## Showcase


Как работает модель на любой картинке из интернета

```
In [13]: model_url = "https://tfhub.dev/google/bit/m-r50x1/ilsvrc2012_classification/
imagenet_module = hub.KerasLayer(model_url)
```

```
In [16]: # Load image (image provided is CC0 licensed)
img_url = "http://images6.fanpop.com/image/photos/34700000/Grey-Elephant-col
image = load_image_from_url(img_url)

# Run model on image
logits = imagenet_module(image)

# Show image and predictions
show_preds(logits, image[0])
```



Как работает модель на данных из датасета. Модель иногда понимает что видит цветок, однако понимания какой именно - нет, много путает с другими объектами.

```
In [17]: train_split = 0.9
num_train = int(train_split * num_examples)
ds_train = ds.take(num_train)
ds_test = ds.skip(num_train)

DATASET_NUM_TRAIN_EXAMPLES = num_examples
```

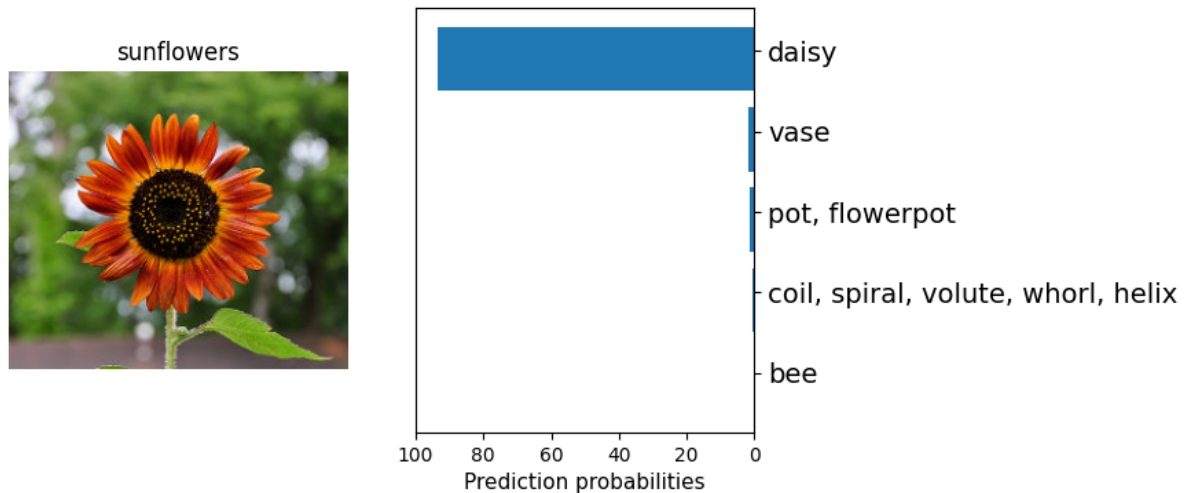
```
In [18]: for features in ds_train.take(1):
    image = features['image']
    image = preprocess_image(image)

    # Run model on image
    logits = imagenet_module(image)
```

```
# Show image and predictions
```

```
show_preds(logits, image[0], correct_flowers_label=features['label'].numpy
```

```
2022-12-26 14:02:57.453063: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```



## Перенос (дообучение) на нашем датасете

### Гиперпараметры и прочее

```
In [19]: # Add new head to the BiT model
```

```
class MyBiTModel(tf.keras.Model):  
    """BiT with a new head."""  
  
    def __init__(self, num_classes, module):  
        super().__init__()  
  
        self.num_classes = num_classes  
        self.head = tf.keras.layers.Dense(num_classes, kernel_initializer='zeros')  
        self.bit_model = module  
  
    def call(self, images):  
        # No need to cut head off since we are using feature extractor model  
        bit_embedding = self.bit_model(images)  
        return self.head(bit_embedding)  
  
model = MyBiTModel(num_classes=NUM_CLASSES, module=module)
```

```
In [20]: IMAGE_SIZE = "\u003C96x96 px" #@param ["<96x96 px", "> 96 x 96 px"]  
DATASET_SIZE = "\u003C20k examples" #@param ["<20k examples", "20k-500k exam
```

```
if IMAGE_SIZE == "<96x96 px":  
    RESIZE_TO = 160  
    CROP_TO = 128  
else:  
    RESIZE_TO = 512
```

```

CROP_TO = 480

if DATASET_SIZE == "<20k examples":
    SCHEDULE_LENGTH = 500
    SCHEDULE_BOUNDARIES = [200, 300, 400]
elif DATASET_SIZE == "20k-500k examples":
    SCHEDULE_LENGTH = 10000
    SCHEDULE_BOUNDARIES = [3000, 6000, 9000]
else:
    SCHEDULE_LENGTH = 20000
    SCHEDULE_BOUNDARIES = [6000, 12000, 18000]

```

```

In [21]: # Preprocessing helper functions

# Create data pipelines for training and testing:
BATCH_SIZE = 512
SCHEDULE_LENGTH = SCHEDULE_LENGTH * 512 / BATCH_SIZE

STEPS_PER_EPOCH = 10

def cast_to_tuple(features):
    return (features['image'], features['label'])

def preprocess_train(features):
    # Apply random crops and horizontal flips for all tasks
    # except those for which cropping or flipping destroys the label semantics
    # (e.g. predict orientation of an object)
    features['image'] = tf.image.random_flip_left_right(features['image'])
    features['image'] = tf.image.resize(features['image'], [RESIZE_TO, RESIZE_TO])
    features['image'] = tf.image.random_crop(features['image'], [CROP_TO, CROP_TO])
    features['image'] = tf.cast(features['image'], tf.float32) / 255.0
    return features

def preprocess_test(features):
    features['image'] = tf.image.resize(features['image'], [RESIZE_TO, RESIZE_TO])
    features['image'] = tf.cast(features['image'], tf.float32) / 255.0
    return features

pipeline_train = (ds_train
    .shuffle(10000)
    .repeat(int(SCHEDULE_LENGTH * BATCH_SIZE / DATASET_NUM_TRAIN))
    .map(preprocess_train, num_parallel_calls=8)
    .batch(BATCH_SIZE)
    .map(cast_to_tuple) # for keras model.fit
    .prefetch(2))

pipeline_test = (ds_test.map(preprocess_test, num_parallel_calls=1)
    .map(cast_to_tuple) # for keras model.fit
    .batch(BATCH_SIZE)
    .prefetch(2))

```

```

In [22]: # Define optimiser and loss

lr = 0.003 * BATCH_SIZE / 512

```

```
# Decay learning rate by a factor of 10 at SCHEDULE_BOUNDARIES.
lr_schedule = tf.keras.optimizers.schedules.PiecewiseConstantDecay(boundaries=[1, 2, 3, 4, 5, 6, 7, 8, 9],
                                                                    values=[1, 0.1, 0.01, 0.001, 0.0001, 0.00001, 0.000001, 0.0000001, 0.00000001])
optimizer = tf.keras.optimizers.SGD(learning_rate=lr_schedule, momentum=0.9)

loss_fn = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True)
```

## Обучение

In [27]: `%load_ext tensorboard`

```
In [32]: # model.compile(optimizer=optimizer,
#                       loss=loss_fn,
#                       metrics=['accuracy'])

# log_dir = "./logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
# tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir)#, h

# # histogram_freq=1 - на маке полностью сломано (расчет гистограммы каждую эпоху)

# # Fine-tune model
# history = model.fit(
#     pipeline_train,
#     batch_size=BATCH_SIZE,
#     steps_per_epoch=STEPS_PER_EPOCH,
#     epochs= 10,
#     validation_data=pipeline_test,
#     callbacks=[tensorboard_callback]
# )
```

Epoch 1/10

2022-12-26 14:40:56.799844: I tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113] Plugin optimizer for device\_type GPU is enabled.

10/10 [=====] - ETA: 0s - loss: 0.0626 - accuracy: 0.9812

2022-12-26 14:41:36.350812: I tensorflow/core/grappler/optimizers/custom\_graph\_optimizer\_registry.cc:113] Plugin optimizer for device\_type GPU is enabled.

WARNING:tensorflow:6 out of the last 6 calls to <function Model.make\_test\_function.<locals>.test\_function at 0x2e6ce8ee0> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce\_retracing=True option that can avoid unnecessary retracing. For (3), please refer to [https://www.tensorflow.org/guide/function#controlling\\_retracing](https://www.tensorflow.org/guide/function#controlling_retracing) and [https://www.tensorflow.org/api\\_docs/python/tf/function](https://www.tensorflow.org/api_docs/python/tf/function) for more details.

WARNING:tensorflow:6 out of the last 6 calls to <function Model.make\_test\_function.<locals>.test\_function at 0x2e6ce8ee0> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce\_retracing=True option that can avoid unnecessary retracing. For (3), please refer to [https://www.tensorflow.org/guide/function#controlling\\_retracing](https://www.tensorflow.org/guide/function#controlling_retracing) and [https://www.tensorflow.org/api\\_docs/python/tf/function](https://www.tensorflow.org/api_docs/python/tf/function) for more details.

```
10/10 [=====] - 59s 5s/step - loss: 0.0626 - accuracy: 0.9812 - val_loss: 0.0389 - val_accuracy: 0.9891
Epoch 2/10
10/10 [=====] - 25s 3s/step - loss: 0.0599 - accuracy: 0.9826 - val_loss: 0.0699 - val_accuracy: 0.9782
Epoch 3/10
10/10 [=====] - 26s 3s/step - loss: 0.0542 - accuracy: 0.9822 - val_loss: 0.0503 - val_accuracy: 0.9891
Epoch 4/10
10/10 [=====] - 24s 2s/step - loss: 0.0512 - accuracy: 0.9863 - val_loss: 0.0231 - val_accuracy: 0.9973
Epoch 5/10
10/10 [=====] - 26s 3s/step - loss: 0.0441 - accuracy: 0.9852 - val_loss: 0.0296 - val_accuracy: 0.9864
Epoch 6/10
10/10 [=====] - 23s 2s/step - loss: 0.0424 - accuracy: 0.9863 - val_loss: 0.0431 - val_accuracy: 0.9837
Epoch 7/10
10/10 [=====] - 24s 2s/step - loss: 0.0371 - accuracy: 0.9873 - val_loss: 0.0158 - val_accuracy: 1.0000
Epoch 8/10
10/10 [=====] - 26s 3s/step - loss: 0.0369 - accuracy: 0.9904 - val_loss: 0.0278 - val_accuracy: 0.9946
Epoch 9/10
10/10 [=====] - 23s 2s/step - loss: 0.0391 - accuracy: 0.9875 - val_loss: 0.0458 - val_accuracy: 0.9864
Epoch 10/10
10/10 [=====] - 22s 2s/step - loss: 0.0317 - accuracy: 0.9906 - val_loss: 0.0469 - val_accuracy: 0.9864
```

In [33]: `%tensorboard --logdir logs/fit`

```
In [36]: # #Save fine-tuned model as SavedModel
# export_module_dir = './tmp/my_saved_bit_model/'
# tf.saved_model.save(model, export_module_dir)
```

```
INFO:tensorflow:Assets written to: ./tmp/my_saved_bit_model/assets
```

```
INFO:tensorflow:Assets written to: ./tmp/my_saved_bit_model/assets
```

## Загрузка из сохраненного стейта и showcase

```
In [37]: saved_module = hub.KerasLayer(export_module_dir, trainable=True)
```

```
2022-12-26 14:49:02.825135: W tensorflow/core/common_runtime/graph_constructor.cc:805] Node 're_lu_48/PartitionedCall' has 1 outputs but the _output_shapes attribute specifies shapes for 2 outputs. Output shapes may be inaccurate.
```

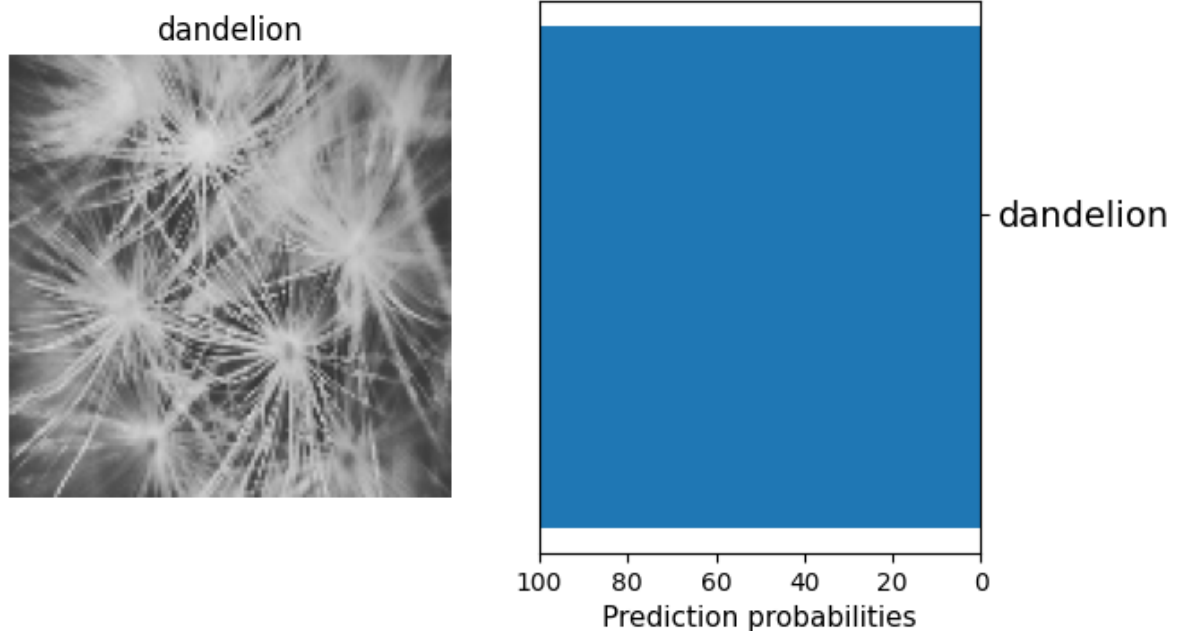
```
2022-12-26 14:49:02.825908: W tensorflow/core/common_runtime/graph_constructor.cc:805] Node 'global_average_pooling2d/PartitionedCall' has 1 outputs but the _output_shapes attribute specifies shapes for 4 outputs. Output shapes may be inaccurate.
```

```
In [38]: for features in ds_train.take(1):
        image = features['image']
        image = preprocess_image(image)
        image = tf.image.resize(image, [CROP_T0, CROP_T0])

        # Run model on image
        logits = saved_module(image)

        # Show image and predictions
        show_preds(logits, image[0], correct_flowers_label=features['label'].numpy)
```

```
2022-12-26 14:49:07.590225: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:113] Plugin optimizer for device_type GPU is enabled.
```



## Preprocessing ESRGAN

[Enhanced Super Resolution GAN](#)

Модель для обработки изображений в супер-разрешение.

## Хелпер-функции

```
In [180... def preprocess_image(image_path):
    """ Loads image from path and preprocesses to make it model ready
    Args:
        image_path: Path to the image file
    """
    hr_image = tf.image.decode_image(tf.io.read_file(image_path))
    # If PNG, remove the alpha channel. The model only supports
    # images with 3 color channels.
    if hr_image.shape[-1] == 4:
        hr_image = hr_image[..., :-1]
    hr_size = (tf.convert_to_tensor(hr_image.shape[:-1]) // 4) * 4
    hr_image = tf.image.crop_to_bounding_box(hr_image, 0, 0, hr_size[0], hr_si
    hr_image = tf.cast(hr_image, tf.float32)
    return tf.expand_dims(hr_image, 0)

def save_image(image, filename):
    """
    Saves unscaled Tensor Images.
    Args:
        image: 3D image tensor. [height, width, channels]
        filename: Name of the file to save.
    """
    if not isinstance(image, Image.Image):
        image = tf.clip_by_value(image, 0, 255)
        image = Image.fromarray(tf.cast(image, tf.uint8).numpy())
    image.save("%s.jpg" % filename)
    print("Saved as %s.jpg" % filename)

def plot_image(image, title=""):
    """
    Plots images from image tensors.
    Args:
        image: 3D image tensor. [height, width, channels].
        title: Title to display in the plot.
    """
    image = np.asarray(image)
    image = tf.clip_by_value(image, 0, 255)
    image = Image.fromarray(tf.cast(image, tf.uint8).numpy())
    plt.imshow(image)
    plt.axis("off")
    plt.title(title)
```

## Применение ESRGAN

```
In [181... ESR_SAVED_MODEL_PATH = "https://tfhub.dev/captain-pool/esrgan-tf2/1"
IMAGE_PATH="image.jpg"
```

Загрузим картинку из датасета и используем готовые веса модели ESRGAN из tfhub



```
In [182... hr_image = preprocess_image(IMAGE_PATH)
plot_image(tf.squeeze(hr_image), title="Original Image")
save_image(tf.squeeze(hr_image), filename="Original Image")
```

Saved as Original Image.jpg

Original Image



```
In [183... model = hub.load(ESR_SAVED_MODEL_PATH)
```

Пропустим это изображение через ESRGAN, время потраченное на одну картинку достаточно мало чтобы обработать все изображения в датасете.

```
In [185... start = time.time()
fake_image = model(hr_image)
fake_image = tf.squeeze(fake_image)
print("Time Taken: %f" % (time.time() - start))
plot_image(tf.squeeze(fake_image), title="Super Resolution")
save_image(tf.squeeze(fake_image), filename="Super Resolution")
```

Time Taken: 0.193167

Saved as Super Resolution.jpg

## Super Resolution



## Выводы

В курсовой был изучен и применен датасет oxford102. На датасете была дообучена модель ResNet50, для улучшения изображений была применена модель ESRGAN, метрики обучения выведены в TensorBoard. Построены графики статистики датасета в Plotly.

Для принта в pdf

```
In [186... !export PATH=/Library/TeX/texbin:$PATH
```

```
In [ ]:
```